

Challenge A:

Overview:

In the humanitarian and international development sector, small to medium-sized non-government organisations (NGOs) deliver a range of impactful programs. These organisations can often see the benefits of adopting AI, but do not have the human or financial resources within the organisation to have dedicated AI technology adoption teams. For many NGOs, ensuring the accuracy of any reported data and the originality of written work, is also key to maintaining trust and reputation with the communities and donors they work with.

Key:

Audience / Stakeholders

Variable to Measure

Constraint

Functional Requirement

Unknowns:

IT resources, cultural appropriateness, sustainability mechanism, harm to third parties, language barriers

Problem: Supporting the adoption of AI across non-governmental organisations (NGOs)

Task: How can small to medium-sized NGOs implement systems and approaches that foster AI adoption while ensuring the authenticity and quality of any analysis or reporting?

Primary Problem:

How might we give resource-constrained small to medium-sized NGOs in the Pacific Islands region lack the human, financial, and IT resources to safely and effectively **adopt AI for reporting with authenticity/accuracy safeguards built directly into that adoption path**. Any adoption path must not compromise the accuracy and originality of their reporting, since that reporting is the foundation of trust with the communities they serve and the donors who fund them.

Key Issues:

1. Might be a bit too vague especially with how many NGOs there are
- **Resources considered**
The brief states outright that these NGOs lack dedicated AI adoption teams or budgets. Any solution assuming ongoing technical maintenance, paid API access, or in-house AI literacy contradicts the premise.
- **Authenticity/originality verification**
The brief flags ensuring "accuracy of any reported data and originality of written work." AI detection and fact-verification tools are **unreliable** and **degrade** fast as generation models improve as a verification layer built today may be obsolete or trivially bypassed within a year.
- **Trust**
Donor and community trust is explicitly named as the stake. If an NGO adopts AI for reporting and it later turns out something was fabricated or subtly wrong (hallucinated statistics, misattributed quotes), the reputational damage likely exceeds whatever efficiency was gained. "Do no harm" as a plausible-sounding but wrong AI-generated report to a donor could do more harm than not adopting AI at all.

- **Connectivity and infrastructure**

Same Pacific geographic context applies. Many field-level NGO staff work in low-connectivity settings. If the solution assumes cloud-based AI tools with reliable internet, it may not function where the reporting/data collection actually happens (in-community), only back at head office. (To counter this Starlink might be a viable option)

- **Capacity and Training**

Small NGOs may have zero staff with AI literacy. This requires hiring new already trained employees or upskilling general employees to be able to use AI technology

- **Who has access**

it's less about community data exposure and more about whose reporting gets AI-assisted. If only well-resourced NGOs or head-office staff can use the tool, it could widen the gap between larger and smaller/grassroots organizations

- **Malicious intent**

Misuse of AI to fabricate or inflate reported outcomes, either by bad actors within an org, or unintentionally through over-trusting AI output.

Potential Solution Ideas:

1. **A reporting / policy assistant** which is made up of two systems. First an LLM drafting tool that helps NGO staff turn field notes, survey data, or raw observations into donor reports and public communications based on checked policies when first setting up the agent. The second system is for fact checking/verification. Every AI-generated claim, statistic, or sentence is automatically tagged with (a) its source document/data point and (b) a simple, human-readable disclosure label shown to both donors and communities (AI declaration) Essentially a website that allows NGO staff to talk to a chatgpt wrapper that has extra guardrails (RAG and rules engine) with extra disclosures about how it might be wrong.

Development Stack

Frontend: web app (works in-browser, minimal install footprint)

Generation: LLM API (online) + small local quantized model (offline fallback)

Retrieval/provenance: Vector store (lightweight, e.g. local embeddings DB) for RAG

Consistency: checkRules engine (regex/numeric matching)

Output: Templated PDF export, serverless deployment architecture

Storage: Local-first data storage with optional sync when connectivity available

Issues with this solution

1. Verification is bounded, not absolute.
2. Human capacity is still required, just less of it
3. Cultural appropriateness of the disclosure format is untested might need more research on this
4. Ongoing cost if API-based / Maintenance feasibility